

EXPERIMENT - 9B

Aim: Visualization using Power BI and Performing EDA, logistic and linear regression using Apache Spark.

About Dataset:

Link to our dataset: <https://www.kaggle.com/datasets/divyansh22/flight-delay-prediction>

This is the first part of flight delay prediction i.e. for the month of January. This data is collected from the Bureau of Transportation Statistics, Govt. of the USA. This data is open-sourced under U.S. Govt. Works. This dataset contains all the flights in the month of January 2019 and January 2020.

There are more than 400,000 flights in the month of January itself throughout the United States. The features were manually chosen to do a primary time series analysis. There are several other features available on their website.

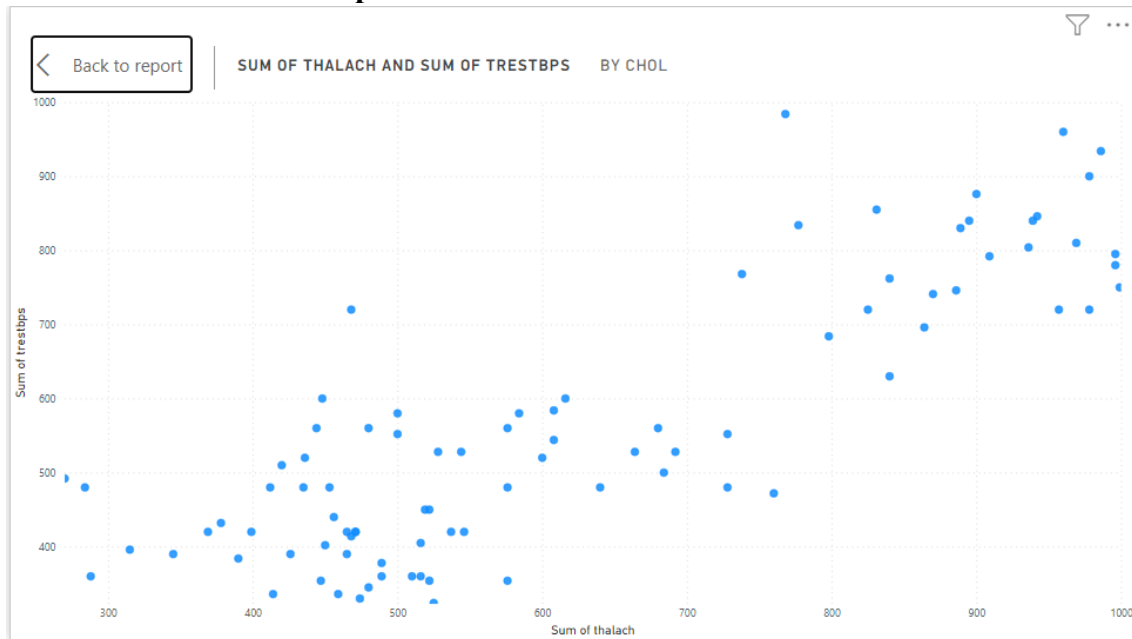
Theory:

Spark is an Apache project advertised as “lightning fast cluster computing”. It has a thriving open-source community and is the most active Apache project at the moment. Spark provides a faster and more general data processing platform. Spark lets you run programs up to 100x faster in memory, or 10x faster on disk, than Hadoop. Last year, Spark took over Hadoop by completing the 100 TB Daytona GraySort contest 3x faster on one tenth the number of machines and it also became the fastest open source engine for sorting a petabyte. Spark also makes it possible to write code more quickly as you have over 80 high-level operators at your disposal. Another important aspect when learning how to use Apache Spark is the interactive shell (REPL) which it provides out-of-the box. Using REPL, one can test the outcome of each line of code without first needing to code and execute the entire job. The path to working code is thus much shorter and ad-hoc data analysis is made possible.

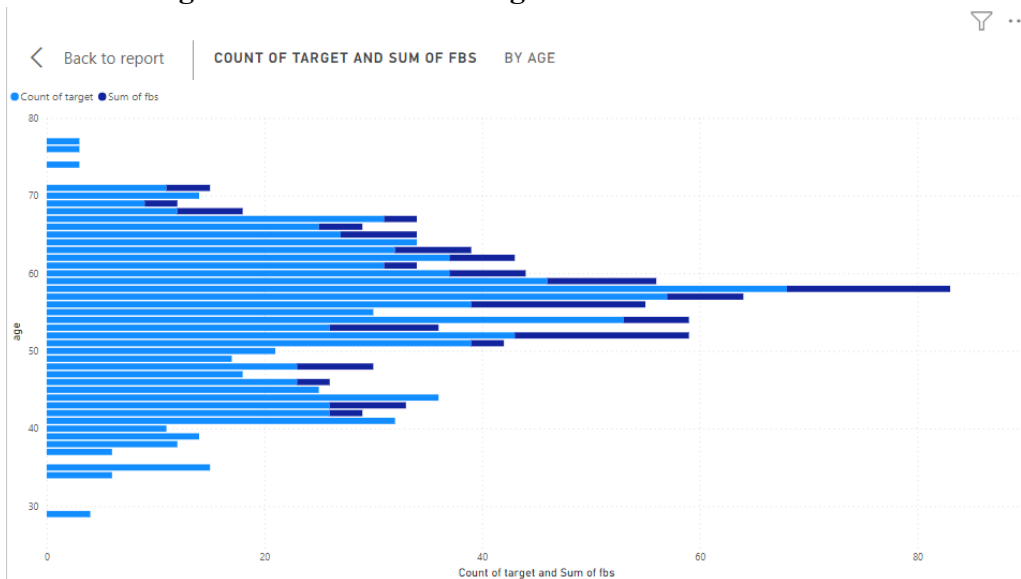
Power BI is a data visualization platform that is primarily used for business intelligence. Power BI's dashboard, which is intended for use by business professionals with varying levels of data knowledge, is capable of reporting and visualizing data in a variety of formats, including graphs, maps, charts, scatter plots, and more. Power BI is made up of several interconnected applications, including Power BI Desktop, Pro, Premium, Mobile, Embedded, and Report Server. While some of these apps are free to use, paid subscriptions to the pro and premium versions offer enhanced analytics capabilities.

Implementation:**Power BI**


1) Scatterplot

Sum of thalach and trestbps vs chol


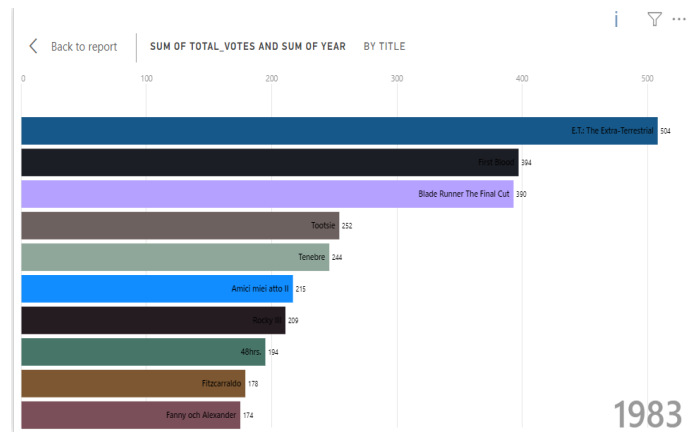
2) Bar graph

Count of Target and Sum of Fbs vs Age

Count of Target by age: It shows the percentage of each age which are likely to get heart disease.




4) Animated visualization



5) Waterfall graph

Sum of trestbps vs thal and age

Apache Spark:-

Step 1: Loading the dataset

Cmd 1

```
1 from pyspark.sql import SparkSession
```

Command took 2.22 seconds -- by 2020.mayuri.yerande@ves.ac.in at 4/5/2023, 4:44:46 PM on My Cluster 5

Cmd 2

```
1 spark = SparkSession.builder.appName('flight').getOrCreate()
```

Command took 0.36 seconds -- by 2020.mayuri.yerande@ves.ac.in at 4/5/2023, 4:44:46 PM on My Cluster 5

Cmd 3

```
1 df = spark.read.format("csv").option("header", "true").load("dbfs:/FileStore/shared_uploads/2020.mayuri.yerande@ves.ac.in/Jan_2020_ontime-2.csv")
```

► (1) Spark Jobs

► df: pyspark.sql.dataframe.DataFrame = [DAY_OF_MONTH: string, DAY_OF_WEEK: string ... 20 more fields]

Command took 18.60 seconds -- by 2020.mayuri.yerande@ves.ac.in at 4/5/2023, 4:44:48 PM on My Cluster 5

```
1 df.printSchema()
```

```
root
|-- DAY_OF_MONTH: string (nullable = true)
|-- DAY_OF_WEEK: string (nullable = true)
|-- OP_UNIQUE_CARRIER: string (nullable = true)
|-- OP_CARRIER_AIRLINE_ID: string (nullable = true)
|-- OP_CARRIER: string (nullable = true)
|-- TAIL_NUM: string (nullable = true)
|-- OP_CARRIER_FL_NUM: string (nullable = true)
|-- ORIGIN_AIRPORT_ID: string (nullable = true)
|-- ORIGIN_AIRPORT_SEQ_ID: string (nullable = true)
|-- ORIGIN: string (nullable = true)
|-- DEST_AIRPORT_ID: string (nullable = true)
|-- DEST_AIRPORT_SEQ_ID: string (nullable = true)
|-- DEST: string (nullable = true)
|-- DEP_TIME: string (nullable = true)
|-- DEP_DEL15: string (nullable = true)
|-- DEP_TIME_BLK: string (nullable = true)
|-- ARR_TIME: string (nullable = true)
|-- ARR_DEL15: string (nullable = true)
|-- CANCELLED: string (nullable = true)
|-- DIVERTED: string (nullable = true)
```

```
Command took 0.23 seconds -- by 2020.mayuri.yerande@ves.ac.in at 4/5/2023, 4:48:32 PM on My Cluster 5
```

```
1 df.groupBy('DEST').count().show()
```

▶ (2) Spark Jobs

```
+-----+-----+
|DEST|count|
+-----+-----+
|BGM|61|
|INL|54|
|PSE|68|
|MSY|4708|
|PPG|9|
|GEG|1111|
|DRT|59|
|BUR|2775|
|SNA|3296|
|GRB|404|
|GTF|159|
|IDA|139|
```

[illegible]

```
1 #converting needed columns into inter type
2 from pyspark.sql import functions as f
3 from pyspark.sql.types import IntegerType
4
5 numeric_columns = ['DAY_OF_MONTH',
6 'DAY_OF_WEEK',
7 'OP_UNIQUE_CARRIER',
8 'OP_CARRIER_AIRLINE_ID',
9 'OP_CARRIER',
10 'TAIL_NUM',
11 'OP_CARRIER_FL_NUM', 'ORIGIN_AIRPORT_ID', 'ORIGIN_AIRPORT_SEQ_ID']
12 for column in numeric_columns:
13     df = df.withColumn(column, f.col(column).cast(IntegerType()))
14 df.printSchema()
```

```
df: pyspark.sql.dataframe.DataFrame = [DAY_OF_MONTH: integer, DAY_OF_WEEK: integer ... 20 more fields]
```

```
root
|-- DAY_OF_MONTH: integer (nullable = true)
|-- DAY_OF_WEEK: integer (nullable = true)
|-- OP_UNIQUE_CARRIER: integer (nullable = true)
|-- OP_CARRIER_AIRLINE_ID: integer (nullable = true)
|-- OP_CARRIER: integer (nullable = true)
|-- TAIL_NUM: integer (nullable = true)
|-- OP_CARRIER_FL_NUM: integer (nullable = true)
|-- ORIGIN_AIRPORT_ID: integer (nullable = true)
|-- ORIGIN_AIRPORT_SEQ_ID: integer (nullable = true)
|-- ORIGIN: string (nullable = true)
|-- DEST_AIRPORT_ID: string (nullable = true)
```

Step 4: Statistical measures of data

```

1  #mean
2  from pyspark.sql.functions import mean, col
3  df_stats = df.select(
4      mean(col('DAY_OF_MONTH')).alias('Mean DAY_OF_MONTH'),
5      mean(col('ARR_TIME')).alias('Mean ARR_TIME'),
6      mean(col('DEP_TIME')).alias('DEP_TIME'),
7
8
9  ).collect()
10 for i in df_stats:
11     row = i.asDict()
12     for k in row:
13         print(k, " - ", row[k])

```

► (2) Spark Jobs

```

Mean DAY_OF_MONTH - 16.014354256058326
Mean ARR_TIME - 1477.9689240359771
DEP_TIME - 1331.512559057871

```

```

1  print("Median DAY_OF_MONTH - ",df.approxQuantile("DAY_OF_MONTH", [0.5], 0.25))

```

► (1) Spark Jobs

```

Median DAY_OF_MONTH - [9.0]

```

Command took 4.17 seconds -- by 2020.mayuri.yerande@ves.ac.in at 4/5/2023, 5:05:19 PM on My Cluster 5

Step 5: Correlation analysis - Pearson

```

1  import pandas as pd
2  from pyspark.mllib.stat import Statistics
3
4  features = df.select(numeric_columns).rdd.map(lambda row: row[0:])
5
6  corr_mat=Statistics.corr(features, method="pearson")
7
8  corr_df = pd.DataFrame(corr_mat,index=numeric_columns, columns=numeric_columns)
9  corr_df

```

► (4) Spark Jobs

	DAY_OF_MONTH	DAY_OF_WEEK	OP_UNIQUE_CARRIER	OP_CARRIER_AIRLINE_ID	OP_CARRIER	TAIL_NUM	OP_CARRIER_FL_NUM	ORIGIN_AIRPORT_ID	ORIGIN_AIRPORT_SEQ_ID
DAY_OF_MONTH	1.000000	-0.053947	NaN	-0.001769	NaN	NaN	-0.020052	-0.004940	
DAY_OF_WEEK	-0.053947	1.000000	NaN	0.007127	NaN	NaN	0.056448	0.005209	
OP_UNIQUE_CARRIER	NaN	NaN	1.0	NaN	NaN	NaN	NaN	NaN	
OP_CARRIER_AIRLINE_ID	-0.001769	0.007127	NaN	1.000000	NaN	NaN	0.425973	-0.046500	
OP_CARRIER	NaN	NaN	NaN	NaN	1.0	NaN	NaN	NaN	
TAIL_NUM	NaN	NaN	NaN	NaN	NaN	1.0	NaN	NaN	
OP_CARRIER_FL_NUM	-0.020052	0.056448	NaN	0.425973	NaN	NaN	1.000000	-0.052829	
ORIGIN_AIRPORT_ID	-0.004940	0.005209	NaN	-0.046500	NaN	NaN	-0.052829	1.000000	
ORIGIN_AIRPORT_SEQ_ID	-0.004940	0.005209	NaN	-0.046500	NaN	NaN	-0.052830	1.000000	

1

2

3

4

5

6

7

8

9

10

```
# get a boolean dataframe where true means that a pair of variables is highly correlated
highly_correlated_df = (abs(corr_df) > .5) & (corr_df < 1.0)

# get the names of the variables so we can use them to slice the dataframe
correlated_vars_index = (highly_correlated_df==True).any()
correlated_var_names = correlated_vars_index[correlated_vars_index==True].index

# slice it
highly_correlated_df.loc[correlated_var_names,correlated_var_names]
```

	ORIGIN_AIRPORT_ID	ORIGIN_AIRPORT_SEQ_ID
ORIGIN_AIRPORT_ID	False	True
ORIGIN_AIRPORT_SEQ_ID	True	False

Command took 0.15 seconds -- by 2020.mayuri.yerande@ves.ac.in at 4/5/2023, 5:06:59 PM on My Cluster 5

1

2

3

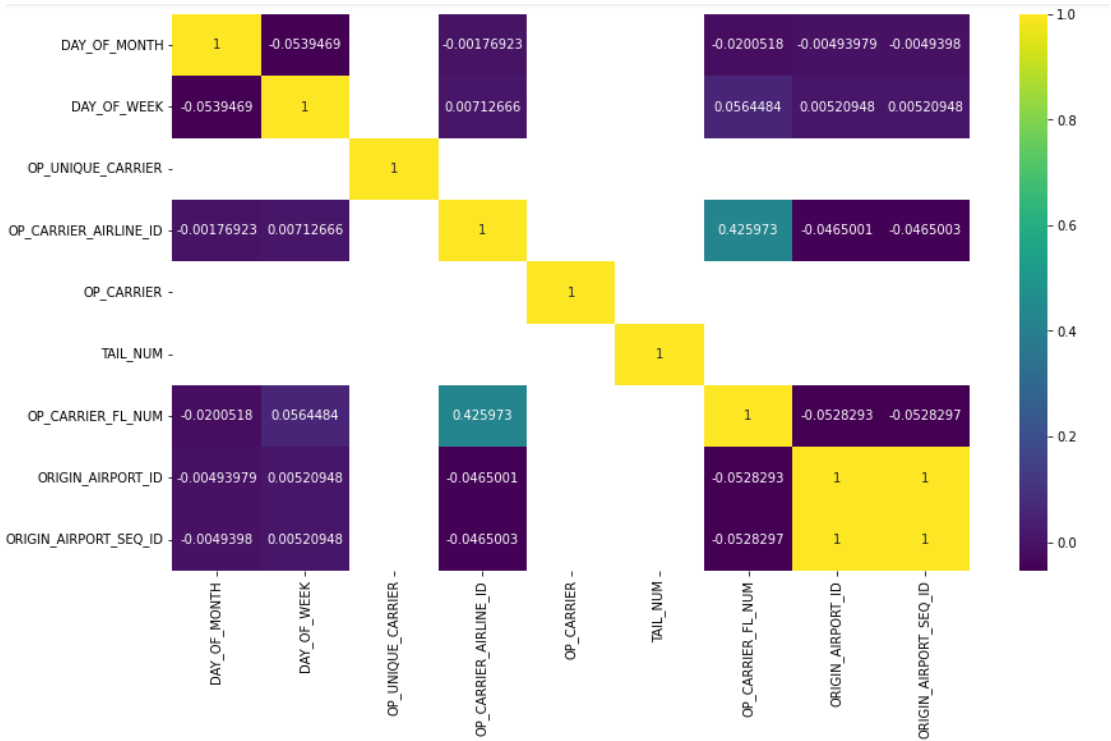
4

5

6

```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(14,8))
sns.heatmap(corr_df, annot=True, fmt="g", cmap='viridis')
```




Step 6: Visualization

```
1 # BOXPLOT
2 x = df.select('DAY_OF_MONTH').toPandas()
3 plt.figure(figsize=(10,5))
4 sns.boxplot('DAY_OF_MONTH',data=x)
```

► (1) Spark Jobs

/databricks/python/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning: on 0.12, the only valid positional argument will be `data`, and passing other arguments will result in a FutureWarning or DeprecationWarning.


warnings.warn(
Out[27]: <AxesSubplot: xlabel='DAY_OF_MONTH'>



```
1 plot_df = df.toPandas()
2 labels=plot_df['DAY_OF_MONTH'].value_counts().keys()
3 size=plot_df['DAY_OF_MONTH'].value_counts()
4 plt.figure(figsize=(15,10))
5 plt.pie(size,labels=labels,autopct='%1.1f%%',)
6 plt.title("Flight taken for day of month")
```

► (1) Spark Jobs


Out[42]: Text(0.5, 1.0, 'Flight taken for day of month')



```
1
2 # VIOLIN PLOT
3 x = df.select('DAY_OF_MONTH').toPandas()
4 plt.figure(figsize=(10,5))
5 sns.violinplot('DAY_OF_MONTH',data=x)
```

► (1) Spark Jobs


Out[23]: <AxesSubplot: xlabel='DAY_OF_MONTH'>



```
1 plot_df = df.toPandas()
2 plt.hist(plot_df["DAY_OF_MONTH"],bins=50)
3 plt.xticks(rotation=90)
4 plt.ylabel("no of units sold")
5 plt.title("No of units sold per Day of Month")
```

► (1) Spark Jobs

Out[43]: Text(0.5, 1.0, 'No of units sold per Day of Month')



Linear Regression using Apache Spark

Dataset: https://github.com/LeondraJames/Hyundai-Cruise-Ship-Crew-Prediction/blob/master/cruise_ship_info.csv

- Import library

```
1 from pyspark.sql import SparkSession
```

```
1 from pyspark.ml.linalg import Vectors
2 from pyspark.ml.feature import VectorAssembler
```

- Create a new SparkSession using the builder pattern.

```
1 spark = SparkSession.builder.appName('cruise').getOrCreate()
```

- Read the dataset

```
1 df1 = spark.read.format("csv").option("header",
2 "true").load("dbfs:/FileStore/shared_uploads/2020.nandana.nair@ves.ac.in/cruise_ship_info.csv", inferSchema=True,
               header=True)
```

- Printing the schema of dataset

```
1 df1.printSchema()
```

```
root
|-- Ship_name: string (nullable = true)
|-- Cruise_line: string (nullable = true)
|-- Age: integer (nullable = true)
|-- Tonnage: double (nullable = true)
|-- passengers: double (nullable = true)
|-- length: double (nullable = true)
|-- cabins: double (nullable = true)
|-- passenger_density: double (nullable = true)
|-- crew: double (nullable = true)
```

- Understanding the contents of dataset

1 df1.show()

(1) Spark Jobs

	Ship_name	Cruise_line	Age	Tonnage	passengers	length	cabins	passenger_density	crew
	Journey	Azamara	6	30.276999999999997	6.94	5.94	3.55	42.64	3.55
	Quest	Azamara	6	30.276999999999997	6.94	5.94	3.55	42.64	3.55
	Celebration	Carnival	26	47.262	14.86	7.22	7.43	31.8	6.7
	Conquest	Carnival	11	110.0	29.74	9.53	14.88	36.99	19.1
	Destiny	Carnival	17	101.353	26.42	8.92	13.21	38.36	10.0
	Ecstasy	Carnival	22	70.367	20.52	8.55	10.2	34.29	9.2
	Elation	Carnival	15	70.367	20.52	8.55	10.2	34.29	9.2
	Fantasy	Carnival	23	70.367	20.56	8.55	10.22	34.23	9.2
	Fascination	Carnival	19	70.367	20.52	8.55	10.2	34.29	9.2
	Freedom	Carnival	6	110.23899999999999	37.0	9.51	14.87	29.79	11.5
	Glory	Carnival	10	110.0	29.74	9.51	14.87	36.99	11.6
	Holiday	Carnival	28	46.052	14.52	7.27	7.26	31.72	6.6
	Imagination	Carnival	18	70.367	20.52	8.55	10.2	34.29	9.2
	Inspiration	Carnival	17	70.367	20.52	8.55	10.2	34.29	9.2
	Legend	Carnival	11	86.0	21.24	9.63	10.62	40.49	9.3
	Liberty*	Carnival	8	110.0	29.74	9.51	14.87	36.99	11.6
	Miracle	Carnival	9	88.5	21.24	9.63	10.62	41.67	10.3
	Paradise	Carnival	15	70.367	20.52	8.55	10.2	34.29	9.2

1 df1.describe().show()

(2) Spark Jobs

	summary	Ship_name	Cruise_line	Age	Tonnage	passengers	length	cabins	passenger_density	crew
	count	158	158	158	158	158	158	158	158	158
	mean	Infinity	null	15.689873417721518	71.28467088607599	18.45740506329114	8.130632911392404	8.830000000000005	39.90094936708861	7.794177215189873
	stddev	null	null	7.615691058751413	37.229540025907866	9.677094775143416	1.793473548054825	4.4714172221480615	8.63921711391542	3.503486564627034
	min	Adventure	Azamara	4	2.329	0.66	2.79	0.33	17.7	0.59
	max	Zuiderdam	Windstar	48	220.0	54.0	11.82	27.0	71.43	21.0

- Count the number of occurrences of each unique value in this column

```
1 df1.groupBy('Cruise_line').count().show()
```

▶ (2) Spark Jobs

```
+-----+-----+
| Cruise_line | count |
+-----+-----+
| Costa       | 11    |
| P&O         | 6     |
| Cunard      | 3     |
| Regent_Seven_Seas | 5    |
| MSC         | 8     |
| Carnival    | 22    |
| Crystal     | 2     |
| Orient      | 1     |
| Princess    | 17    |
| Silversea   | 4     |
| Seabourn    | 3     |
| Holland_American | 14   |
| Windstar    | 3     |
| Disney      | 2     |
| Norwegian   | 13    |
| Oceania     | 3     |
| Azamara     | 2     |
| Celebrity   | 10    |
```

- Convert the categorical column "Cruise_line" in the DataFrame df1 to a numerical index column "cruise_cat" in a new DataFrame indexed

```
1 from pyspark.ml.feature import StringIndexer
2 indexer = StringIndexer(inputCol="Cruise_line", outputCol="cruise_cat")
3 indexed = indexer.fit(df1).transform(df1)
4 indexed.head(5)
```

- Printing new dataframe

```
1 indexed.columns
```

```
Out[54]: ['Ship_name',
'Cruise_line',
'Age',
'Tonnage',
'passengers',
'length',
'cabins',
'passenger_density',
'crew',
'cruise_cat']
```

- Create a VectorAssembler object from the pyspark.ml.feature module that combines multiple input columns in a DataFrame into a single output column.

```
1 assembler = VectorAssembler(
2     inputCols=['Age',
3               'Tonnage',
4               'passengers',
5               'length',
6               'cabins',
7               'passenger_density',
8               'cruise_cat'],
9     outputCol="features")
```

```
1 output = assembler.transform(indexed)
```

```
1 output.select("features", "crew").show()
```

▶ (1) Spark Jobs

```
+-----+-----+
|          features|crew|
+-----+-----+
|[6.0,30.276999999...|3.55|
|[6.0,30.276999999...|3.55|
|[26.0,47.262,14.8...| 6.7|
|[11.0,110.0,29.74...|19.1|
|[17.0,101.353,26...|10.0|
|[22.0,70.367,20.5...| 9.2|
|[15.0,70.367,20.5...| 9.2|
|[23.0,70.367,20.5...| 9.2|
|[19.0,70.367,20.5...| 9.2|
|[6.0,110.23899999...|11.5|
|[10.0,110.0,29.74...|11.6|
|[28.0,46.052,14.5...| 6.6|
|[18.0,70.367,20.5...| 9.2|
|[17.0,70.367,20.5...| 9.2|
|[11.0,86.0,21.24,...| 9.3|
|[8.0,110.0,29.74,...|11.6|
|[9.0,88.5,21.24,9...|10.3|
|[15.0,70.367,20.5...| 9.2|
```

- Select the "features" column and the "crew" column from the output DataFrame

```
1 final_data = output.select("features", "crew")
```

- Splitting the dataset into train and test dataset

```
1 train_data, test_data = final_data.randomSplit([0.7,0.3])
```

- Model creation

```
1 from pyspark.ml.regression import LinearRegression
2 # Create a Linear Regression Model object
3 lr = LinearRegression(labelCol='crew')
```

- Fitting the model

```
1 # Fit the model to the data and call this model lrModel
2 lrModel = lr.fit(train_data)
```

- Printing coefficients and intercept

```

1 # Print the coefficients and intercept for linear regression
2 print("Coefficients: {} Intercept: {}".format(lrModel.coef, lrModel.intercept))

```

Coefficients: [-0.007907524274939418, 0.010549738042022238, -0.17312350509313318, 0.38909053958896234, 0.9215589936576951, -0.007093471571446334, 0.05534458494471553] Intercept: -0.8749465017603469

```

1 test_results = lrModel.evaluate(test_data)

```

```

1 print("RMSE: {}".format(test_results.rootMeanSquaredError))
2 print("MSE: {}".format(test_results.meanSquaredError))
3 print("R2: {}".format(test_results.r2))

```

RMSE: 0.6588197285198434
MSE: 0.43404343468696016
R2: 0.9641523051695818

- Performing correlation

```

1 # R2 of 0.86 is pretty good, let's check the data a little closer
2 from pyspark.sql.functions import corr

```

```

1 df1.select(corr('crew', 'passengers')).show()

```

▶ (2) Spark Jobs

```

+-----+
|corr(crew, passengers)|
+-----+
|      0.9152341306065384|
+-----+

```

Command took 0.76 seconds -- by 2020.nandana.nair@ves.ac.in at 2/10/2023, 3:54:58 PM on Session

Cmd 26

```

1 df1.select(corr('crew', 'cabins')).show()

```

▶ (2) Spark Jobs

```

+-----+
|corr(crew, cabins)|
+-----+
|0.9508226063578497|
+-----+

```

Logistic regression using Apache spark

Dataset: https://github.com/SkalskiP/pySpark_Tutorial/blob/master/Sekcja_12_Logistic_Regression/new_customers.csv

- Importing pyspark

```

1 from pyspark.sql import SparkSession

```

```

1 spark = SparkSession.builder.appName('logregconsult').getOrCreate()

```

- Printing the schema for the dataset.

```
1 data.printSchema()
```

```
root
 |-- Names: string (nullable = true)
 |-- Age: double (nullable = true)
 |-- Total_Purchase: double (nullable = true)
 |-- Account_Manager: integer (nullable = true)
 |-- Years: double (nullable = true)
 |-- Num_Sites: double (nullable = true)
 |-- Onboard_date: timestamp (nullable = true)
 |-- Location: string (nullable = true)
 |-- Company: string (nullable = true)
 |-- Churn: integer (nullable = true)
```

- Computes and displays descriptive statistics of a DataFrame

```
1 data.describe().show()
```

▶ (2) Spark Jobs

summary	Names	Age	Total_Purchase	Account_Manager	Years	Num_Sites	Lo
cation	Company	Churn					
count	900	900	900	900	900	900	
900	900	900					
mean	null 41.81666666666667	10062.824033333334	0.4811111111111111	5.273155555555555	8.587777777777777		
null	null 0.16666666666666666						
stddev	null 6.127560416916251	2408.644531858096	0.4999208935073339	1.274449013194616	1.7648355920350969		
null	null 0.3728852122772358						
min	Aaron King	22.0	100.0	0	1.0	3.0 00103	Jeffrey
Cre...	Abbott-Thompson		0				
max	Zachary Walsh	65.0	18026.01	1	9.15	14.0 Unit 9800	Box
287...	Zuniga, Clark and...		1				

```
1 data.columns
```


```
Out[34]: ['Names',
 'Age',
 'Total_Purchase',
 'Account_Manager',
 'Years',
 'Num_Sites',
 'Onboard_date',
 'Location',
 'Company',
 'Churn']
```

- The VectorAssembler is a transformer in PySpark that is used to combine a given list of columns into a single vector column.

```
1 from pyspark.ml.feature import VectorAssembler
```

```
1 assembler = VectorAssembler(inputCols=['Age',
2   'Total_Purchase',
3   'Account_Manager',
4   'Years',
5   'Num_Sites'],outputCol='features')
```

```
1 output = assembler.transform(data)
```

▶  output: pyspark.sql.dataframe.DataFrame = [Names: string, Age: double ... 9 more fields]

- This line of code creates a new DataFrame final_data by selecting two columns from the output DataFrame: the features column and the churn column.

```
1 final_data = output.select('features', 'churn')
```

▶  final_data: pyspark.sql.dataframe.DataFrame = [features: udt, churn: integer]

- Splitting the dataset into train and test set

```
1 train_churn, test_churn = final_data.randomSplit([0.7, 0.3])
```

▶  train_churn: pyspark.sql.dataframe.DataFrame = [features: udt, churn: integer]

▶  test_churn: pyspark.sql.dataframe.DataFrame = [features: udt, churn: integer]

- Importing LogisticRegression library

```
1 from pyspark.ml.classification import LogisticRegression
```

- This line of code initializes a new LogisticRegression estimator in PySpark with the label column set to churn.

```
1 lr_churn = LogisticRegression(labelCol='churn')
```

- Fitting the train set in the model.

```
1 fitted_churn_model = lr_churn.fit(train_churn)
```

```
1 training_sum = fitted_churn_model.summary
```

```
1 training_sum.predictions.describe().show()
```

► (2) Spark Jobs

summary	churn	prediction
count	630	630
mean	0.173015873015873	0.1523809523809524
stddev	0.37856156048067147	0.3596753273871411
min	0.0	0.0
max	1.0	1.0

```
1 from pyspark.ml.evaluation import BinaryClassificationEvaluator
```

```
1 pred_and_labels = fitted_churn_model.evaluate(test_churn)
```

- Display the predictions made by the Logistic Regression model.

```
1 pred_and_labels.predictions.show()
```

► (1) Spark Jobs

features	churn	rawPrediction	probability	prediction
[26.0,8787.39,1.0...]	1	[-0.2029268544860...]	[0.44944166386318...]	1.0
[27.0,8628.8,1.0,...]	0	[5.64245690587166...]	[0.99646836658134...]	0.0
[28.0,8670.98,0.0...]	0	[8.50981099862742...]	[0.99979855868510...]	0.0
[28.0,9090.43,1.0...]	0	[1.02136548408360...]	[0.73523849385903...]	0.0
[29.0,12711.15,0.0...]	0	[5.91023747386637...]	[0.99729578969188...]	0.0
[29.0,13240.01,1.0...]	0	[6.97159023565538...]	[0.99906271949514...]	0.0
[30.0,6744.87,0.0...]	0	[3.74110328125562...]	[0.97682205409108...]	0.0
[30.0,10183.98,1.0...]	0	[2.59020233148786...]	[0.93022835030179...]	0.0
[31.0,8829.83,1.0...]	0	[4.40191565403454...]	[0.98789449567262...]	0.0
[31.0,10058.87,1.0...]	0	[4.85040392638843...]	[0.99223554245866...]	0.0
[32.0,7896.65,0.0...]	0	[4.15406075598772...]	[0.98454216537319...]	0.0
[32.0,10716.75,0.0...]	0	[4.81123885322278...]	[0.99192791682664...]	0.0
[32.0,13630.93,0.0...]	0	[1.99417782817638...]	[0.88018443073451...]	0.0
[33.0,8556.73,0.0...]	0	[3.91524217348165...]	[0.98045394453372...]	0.0
[33.0,10306.21,1.0...]	0	[1.94600468809309...]	[0.87501033984067...]	0.0
[33.0,13157.08,1.0...]	0	[0.84511283591903...]	[0.69954094443479...]	0.0
[34.0,5447.16,1.0...]	0	[3.31603538178110...]	[0.96497484046300...]	0.0
[34.0,6131.92,0.0...]	0	[4.10531645935957...]	[0.98378254035572...]	0.0

CONCLUSION: Thus we performed Visualization using Power BI and used its various features. We performed EDA, logistic and linear regression on Apache Spark and did some visualization. Thus we successfully learnt about power BI and Apache Spark